
Operational notes

Document updated on **February 2, 2022**.

The following colors are **not** part of the final product, but serve as highlights in the editing/review process:

- text that needs attention from the Subject Matter Experts: Mirco, Anna,& Jan
- terms that have not yet been defined in the book
- text that needs advice from the communications/marketing team: Aaron & Shane
- text that needs to be completed or otherwise edited (by Sylvia)

NB: This PDF only includes the Circuit Compilers chapter

Todo list

zero-knowledge proofs	12
played with	12
finite field	12
elliptic curve	12
Update reference when content is finalized	12
methatical	12
numerical	12
a list of additional exercises	13
think about them	13
add some more informal explanation of absolute value	14
We haven't really talked about what a ring is at this point	14
What's the significance of this distinction?	15
reverse	15
Turing machine	15
polynomial time	15
sub-exponentially, with $\mathcal{O}((1 + \varepsilon)^n)$ and some $\varepsilon > 0$	15
Add text	16
\mathbb{Q} of fractions	16
Division in the usual sense is not defined for integers	16
Add more explanation of how this works	17
pseudocode	18
modular arithmetics	18
actual division	18
multiplicative inverses	18
factional numbers	18
exponentiation function	20
See XXX	20
once they accept that this is a new kind of calculations, its actually not that hard	20
perform Euclidean division on them	20
This Sage snippet should be described in more detail.	21
prime fields	23
residue class rings	23
Algorithm sometimes floated to the next page, check this for final version	23
Add a number and title to the tables	25
(-1) should be $(-a)$?	26
we have	28
rephrase	32
subtrahend	33
minuend	33

■ what does this mean?	37
■ add reference	114
■ add reference	116
■ add reference	116
■ add reference	116
■ add reference	116
■ add reference	116
■ add reference	117
■ add reference	117
■ add reference	117
■ add references	117
■ add reference	117
■ add reference	118
■ add reference	118
■ add reference	119
■ add reference	119
■ add reference	119
■ add reference	120
■ add reference	120
■ add reference	121
■ add reference	122
■ add reference	122
■ add reference	123
■ add reference	123
■ add reference	125
■ add reference	125
■ add reference	125
■ add reference	125
■ add reference	125
■ add reference	125
■ add reference	127
■ add reference	127
■ add reference	127
■ add reference	128
■ add reference	128
■ add reference	129
■ add reference	131
■ add reference	131
■ add reference	132
■ add reference	132
■ add reference	132
■ add reference	133
■ add reference	133
■ add reference	133
■ add reference	134
■ add reference	135
■ add reference	135
■ add reference	138

■ add reference	139
■ add reference	139
■ add reference	139
■ add reference	139
■ add reference	139
■ add reference	139
■ add reference	139
■ add reference	142
■ add reference	142
■ add reference	142
■ add reference	142

MoonMath manual

TechnoBob and the Least Scruples crew

February 2, 2022

Contents

1	Introduction	5
1.1	Target audience	5
1.2	The Zoo of Zero-Knowledge Proofs	6
	To Do List	8
	Points to cover while writing	8
2	Preliminaries	9
2.1	Preface and Acknowledgements	9
2.2	Purpose of the book	9
2.3	How to read this book	10
2.4	Cryptological Systems	10
2.5	SNARKS	10
2.6	complexity theory	10
2.6.1	Runtime complexity	10
2.7	Software Used in This Book	11
2.7.1	Sagemath	11
3	Arithmetics	12
3.1	Introduction	12
3.1.1	Aims and target audience	12
3.1.2	The structure of this chapter	13
3.2	Integer Arithmetics	13
	Euclidean Division	16
	The Extended Euclidean Algorithm	18
3.3	Modular arithmetic	19
	Congurency	20
	Modular Arithmetics	20
	The Chinese Remainder Theorem	23
	Modular Inverses	26
3.4	Polynomial Arithmetics	29
	Polynomial Arithmetics	33
	Euklidean Division	34
	Prime Factors	36
	Lange interpolation	37
4	Algebra	40
4.1	Groups	40
	Commutative Groups	41
	Finite groups	43

	Generators	43
	The discrete Logarithm problem	43
4.1.1	Cryptographic Groups	44
	The discret logarithm assumption	45
	The decisional Diffi Hellman assumption	47
	The computational Diffi Hellman assumption	47
	Cofactor Clearing	48
4.1.2	Hashing to Groups	48
	Hash functions	48
	Hashing to cyclic groups	50
	Hashing to modular arithmetics	51
	Pederson Hashes	54
	MimC Hashes	55
	Pseudo Random Functions in DDH-A groups	55
4.2	Commutative Rings	55
	Hashing to Commutative Rings	58
4.3	Fields	58
	Prime fields	59
	Square Roots	60
	Exponentiation	62
	Hashing into Prime fields	62
	Extension Fields	62
	Hashing into extension fields	65
4.4	Projective Planes	65
5	Elliptic Curves	68
5.1	Elliptic Curve Arithmetics	68
5.1.1	Short Weierstraß Curves	68
	Affine short Weierstraß form	69
	Affine compressed representation	73
	Affine group law	73
	Scalar multiplication	77
	Projective short Weierstraß form	80
	Projective Group law	81
	Coordinate Transformations	83
5.1.2	Montgomery Curves	83
	Affine Montgomery Form	83
	Affine Montgomery coordinate transformation	85
	Montgomery group law	86
5.1.3	Twisted Edwards Curves	86
	Twisted Edwards Form	87
	Twisted Edwards group law	88
5.2	Elliptic Curves Pairings	89
	Embedding Degrees	89
	Elliptic Curves over extension fields	90
	Full Torsion groups	91
	Torsion-Subgroups	92
	The Weil Pairing	94

5.3	Hashing to Curves	96
	Try and increment hash functions	96
5.4	Constructing elliptic curves	98
	The Trace of Frobenius	99
	The j -invariant	100
	The Complex Multiplication Method	100
	The <i>BLS6_6</i> pen& paper curve	107
	Hashing to the pairing groups	112
6	Statements	114
6.1	Formal Languages	114
	Decision Functions	115
	Instance and Witness	118
	Modularity	120
6.2	Statement Representations	121
6.2.1	Rank-1 Quadratic Constraint Systems	121
	R1CS representation	121
	R1CS Satisfiability	124
	Modularity	125
6.2.2	Algebraic Circuits	126
	Algebraic circuit representation	126
	Circuit Execution	131
	Circuit Satisfiability	132
	Associated Constraint Systems	133
6.2.3	Quadratic Arithmetic Programs	138
	QAP representation	139
	QAP Satisfiability	141
7	Circuit Compiler	144
7.1	A Pen and Paper Language	144
7.1.1	The Grammar	144
7.1.2	The Execution Phases	146
	The Setup Phase	146
	The Proofer Phase	148
7.2	Common Programing concepts	148
7.2.1	Primitive Types	148
	The Basefield type	148
	The Subtraction Constraints System	152
	The Inversion Constraint System	152
	The Division Constraint System	153
	The Boolean Type	155
	The Boolean Constraint System	155
	The AND operator constraint system	156
	The OR operator constraint system	156
	The NOT operator constraint system	157
	Modularity	158
	Arrays	161
	The Unsigned Integer Type	161

	The uN Constraints System	162
	The Unsigned Integer Operators	163
7.2.2	Control Flow	164
	The Conditional Assignment	164
	Loops	166
7.2.3	Binary Field Representations	167
7.2.4	Cryptographic Primitives	169
	Twisted Edwards curves	169
	Twisted Edwards curves constraints	169
	Twisted Edwards curves addition	170
8	Zero Knowledge Protocols	170
8.1	Proof Systems	170
8.2	The "Groth16" Protocol	171
	The Setup Phase	173
	The Proofer Phase	178
	The Verification Phase	181
	Proof Simulation	183
9	Exercises and Solutions	186

Chapter 7

Circuit Compiler

As we have seen in the previous chapter, statements can be formalized as membership or knowledge claims in formal language and both algebraic circuits as well as rank-1 constraints systems are two practically important ways to define those languages.

However both algebraic circuits and rank-1 constraints systems are not ideal from a developers point of view, because they deviate substantially from common programming paradigms. Writing real world applications as circuits and the associated verification in terms of rank-1 constraint systems is at least as troublesome as writing any other low level language like assembler code. To allow for complex statement design it is therefore necessary to have some kind of compiler framework, capable to transform high level languages into arithmetic circuits and associated rank-1 constraint systems.

As we have seen in XXX as well as XXX and XXX, both arithmetic circuits and rank-1 constraint systems have a modularity property by which it is possible to synthesise complex circuits from simple ones. A basic approach taken by many circuit/R1CS compilers is therefore to provide a library of atomic and simple circuits and then define a way to combine those basic building blocks into arbitrary complex systems.

In this chapter we will provide an introduction to basic concepts of so called *circuit compilers* and derive a toy language which we can "compile" in a pen and paper approach into algebraic circuits and their associated rank-1 constraints systems.

We start with a general introduction to our language and then introduce atomic types like booleans, unsigned integers and define the fundamental control flow primitives like the if-then-else conditional and the bounded loop. We will look at basic functionality primitives like elliptic curve cryptography. Primitives like those are often called **gadgets** in the literature.

7.1 A Pen and Paper Language

To explain basic concepts of circuit compiler and their associated high level languages, we derive an informal toy language and associated brain compiler which we name **PAPER (Pen And Paper Execution Rules)**. PAPER allows programmers to define statements in Rust-like pseudo-code. The language is inspired by ZOKRATES and `circom`.

7.1.1 The Grammar

In PAPER any statement is defined as an ordered list of functions, where any function has to be declared in the list before it is called in another function of that list. The last entry in a statement has to be a special function, called `main`. Functions take a list of typed parameters as inputs

and compute a tuple of typed variables as output, where types are special functions that define how to transform that type into another type, ultimately transforming any type into elements of the base field where the circuit is defined over.

Any statement is parameterized over the field that the circuit will be defined on and has additional optional parameters of unsigned type, needed to define the size of array or the counter of bounded loops. The following definition makes the grammar of a statement precise using a command line language like description:

```
statement <Name> {F:<Field> [ , <N_1: unsigned>, ... ] } {
  [fn <Name>([[pub]<Arg>:<Type>, ...]) -> (<Type>, ...)]{
    [let [pub] <Var>:<Type> ;... ]
    [let const <Const>:<Type>=<Value> ;... ]
    Var<==>(fn ([<Arg>|<Const>|<Var>, ...]) | (<Arg>|<Const>|<Var>)) ;
    return (<Var>, ...) ;
  } ;... ]
  fn main([[pub]<Arg>:<Type>, ...]) -> (<Type>, ...){
    [let [pub] <Var>:<Type> ;... ]
    [let const <Const>:<Type>=<Value> ;... ]
    Var<==>(fn ([<Arg>|<Const>|<Var>, ...]) | (<Arg>|<Const>|<Var>)) ;
    return (<Var>, ...) ;
  } ;
}
```

Function arguments and variables are private by default but can be declared as public by the `pub` specifier. Declaring arguments and variables as public always overwrites any previous or conflicting private declaration. Every argument, constant or variable has a type and every type is defined as a function that transforms that type into another type:

```
type <TYPE>( t1 : <TYPE_1>) -> TYPE_2{
  let t2: TYPE_2 <== fn(TYPE_1)
  return t2
}
```

Many real world circuit languages are based on a similar, but of course more sophisticated approach, than PAPER. The purpose of paper is to show basic principles of circuit compilers and their associated high level languages.

Example 124. To get a better understanding of the grammar of PAPER the following is proper high level code that follows the grammar of the PAPER language, assuming that all types in that code have been defined elsewhere.

```
statement MOCK_CODE {F: F_43, N_1 = 1024, N_2 = 8} {
  fn foo(in_1 : F, pub in_2 : TYPE_2) -> F {
    let const c_1 : F = 0 ;
    let const c_2 : TYPE_2 = SOME_VALUE ;
    let pub out_1 : F ;
    out_1<== c_1 ;
    return out_1 ;
  } ;

  fn bar(pub in_1 : F) -> F {
    let out_1 : F ;
    out_1<==foo(in_1);
    return out_1 ;
  } ;
}
```

```

} ;

fn main(in_1 : TYPE_1) -> (F, TYPE_2) {
  let const c_1 : TYPE_1 = SOME_VALUE ;
  let const c_2 : F = 2 ;
  let const c_3 : TYPE_2 = SOME_VALUE ;
  let pub out_1 : F ;
  let out_2 : TYPE_2 ;
  c_1 <== in_1 ;
  out_1 <== foo(c_2) ;
  out_2 <== TYPE_2 ;
  return (out_1, out_2) ;
} ;
}

```

7.1.2 The Execution Phases

In contrast to normal executable programs, programs for circuit compilers have two modes of execution. The first mode, usually called *setup phase*, is executed in order to generate the circuit and its associated rank-1 constraint system, the letter of which is then usually used as input to some zero knowledge proofing system.

The second mode of execution is usually called the *proofer phase* and in this phase a proofer usually computes a valid assignment to the circuit. Depending on the usecase this valid assignment is then either directly used as constructive proof for proper circuit execution or is transferred as input to the proof generation algorithm of some zero knowledge proofing system, where the full size, non hiding constructive proof is processed into a succinct proof with various levels of zero-knowledge.

Modern circuit languages and their associated compilers abstract over those two phases and provide a unified interphase to the developer, who then writes a single program that can be used in both phases.

To give the reader a clear, conceptual distinction between the two phases, PAPER keeps them separated. Code can be brain compiled during the *setup-phase* in a pen and paper approach into visual circuits. Once a circuit is derived it can be executed in a *proofer phase* to generate a valid assignment. The valid assignment is then interpreted as a constructive proof for a knowledge claim in the associated language.

The Setup Phase In PAPER the task of the setup phase is to compile code in the PAPER language into a visual representation of an algebraic circuit. Deriving the circuit from the code in a pen and paper style is what we call *brain compiling*.

Give some statement description that adheres to the correct grammar, we start the circuit development with an empty circuit, compile the main function first and then inductively compile all other functions as they are called during the process.

For every function we compile, we draw a box-node for every argument, every variable and every constant of that function. If the node represents a variable, we label it with that variables name and if it represents a constant, we label it with that constants value. We group arguments into a subgraph labeled "inputs" and return values into a subgraph labeled "outputs". We then group everything into a subgraph and label that subgraph with the functions name.

After this is done, we have to do a consistency and type check for every occurrence of the

assignment operator `<==`. We have to ensure that the expression on the right side of the operator is well defined and that the types of both side match.

Then we compile the right side of every occurrence of the assignment operator `<==`. If the right side is a constant or variable defined in this function, we draw a dotted line from the box-node that represent the left side of `<==`, to the box node that represents the right side of the same operator. If the right side represents an argument of that function we draw a line from the box-node that represent the left side of `<==`, to the box node that represents the right side of the same operator.

If the right side of the `<==` operator is a function, we look into our database, eventually find its associated circuit and draw it. If no circuit is yet associated to that function we repeat the compilation process for that function, drawing edges from the functions argument to its input nodes and edges from the functions output nodes to the nodes on the right side of `<==`.

During that process edge labels are drawn according to the rules from XXX. If the associated variable represents a private value we use the *W* label to indicate a witness and if it represents a public value we use the *I* label to indicate an instance.

Once this is done, we compile all occurring types in a function, by compiling the function of each type. We do this inductively until we reach the type of the base field. Circuits have no notion of types, only of field elements and hence every type needs to be compiled to the field type in a sequence of compilation steps.

Our compilation stops, once we have inductively related all functions by their circuits. The result is a circuit that contains many unnecessary box nodes. In a final optimization step all box nodes that are directly linked to each other are collapsed into a single node and all box nodes that represent the same constants are collapsed into a single node.

Of course PAPER's brain compiler is not properly defined in any formal manner. Its purpose is to highlight important steps that real world compilers undergo in their setup phases.

Example 125 (A trivial Circuit). To give an intuition of how to write and compile circuits in the PAPER language, consider the following statement description:


```
statement trivial_circuit {F:F_13} {
  fn main{F}(in1 : F, pub in2 : F) -> (F,F) {
    let const outc1 : F = 0 ;
    let const incl : F = 7 ;
    let out1 : F ;
    let out2 : F ;
    out1 <== incl;
    out2 <== in1;
    outc1 <== in2;
    return (out1, out2) ;
  }
}
```

To brain compile this statement into an algebraic circuit with PAPER, we start with an empty circuit and evaluate function `main`, which is the only function in this statement.

We draw box-nodes for every argument, every constant and every variable of the function and label them with their names or values, respectively. Then we do a consistency and type check for every `<==` operator in the function. Since the circuit only wires inputs to outputs and all elements have the same type, the check is valid.

Then we evaluate the right side of the assignment operators. Since in our case the right side of each operator is not a function, we draw edges from the box-nodes on the right side to the associated box node on the left side. To label those edges, we use the general rules of algebraic

circuits as defined in XXX. According to those rules every incoming edge of a sink node has a label and every outgoing edge of a source node has a label, if the node is labeled with a variable. Since nodes that represent constants are implicitly assumed to be private and since the public specifier determines if an edge is labeled with W or I , we get the following circuit:



The Proofer Phase In PAPER a so called proofer phase can be executed once the setup phase has generated a circuit image from its associated high level code. It is done by executing the circuit, while assigning proper values to all input nodes of the circuit. However in contrast to most real world compilers, PAPER does not tell the proofer how to find proper input values to a given circuit. Real world programming languages usually provide this data, by computations that are done outside of the circuit.

Example 126. Consider the circuit from example XXX. Valid assignments to this circuit are constructive proofs that the pair of inputs (S_1, S_2) is a point on the tiny-jubjub curve. However the circuit does not provide a way to actually compute proper values for S_1 and S_2 . Any real world system therefore needs an auxiliary computation, that provides those values.

7.2 Common Programing concepts

In this section we cover concepts that appear in almost every programming language and we see how they can be implemented in circuit compilers.

7.2.1 Primitive Types

Primitive data types like booleans, (unsigned) integers, or strings are the most basic building blocks one might expect in every general high level programing language. In order to write statements as computer programs that compile into circuits, it is therefore necessary to implement primitive types as constraints systems and define their associated operations as circuits.

In this section we will look at some common ways to achieve this. After a recapitulation of the atomic type of prime field elements, we start with an implementation of the boolean type and its associated boolean algebra as circuits. After that we define unsigned integers on top of the boolean type, and leave the implementation of signed integers as an exercise to the reader.

It should be noted however that while in common programing languages like C, Go, or Rust primitive data types have a one-to-one correspondence with objects in the computer's memory. This is different for most languages that compile into algebraic circuits. As we will see in the following paragraphs, common primitives like booleans or unsigned integers require many constraints and memory. Primitives different from the underlying field elements can be expensive.

The Basefield type

Since both algebraic circuits and their associated rank-1 constraint systems are defined over a finite field, elements from that field are the atomic informational units in those models. In this

sense field elements $x \in \mathbb{F}$ are for algebraic circuits what bits are for computers.

In PAPER we write F for this type and specify the actual field instance for every statement in curly brackets after the name of that statement. Two functions are associated to this type, which are induced by the *addition* and *multiplication* law in the field F . We write

$$\text{MUL} : F \times F \rightarrow F ; (x, y) \mapsto \text{MUL}(x, y) \quad (7.1)$$

$$\text{ADD} : F \times F \rightarrow F ; (x, y) \mapsto \text{ADD}(x, y) \quad (7.2)$$


Circuit compilers have to compile these functions into the algebraic gates, as explained in XXX. Every other function has to be expressed in terms of them and proper wiring.

To represent addition and multiplication in the PAPER language, we defin the following two functions:

```
fn MUL(x : F, y : F) -> (MUL(x, y) : F) {}
```

```
fn ADD(x : F, y : F) -> (ADD(x, y) : F) {}
```

The compiler then compiles every occurence of the MUL or the ADD function into the following circuits:



Example 127 (Basic gates). To give an intuition of how a real world compiler might transform addition and multiplication in algebraic expressions into a circuit, consider the following PAPER statement:

```
statement basic_ops {F:F_13} {
  fn main(in_1 : F, pub in_2 : F) -> (out_1:F, out_2:F){
    out_1 <== MUL(in_1, in_2) ;
    out_2 <== ADD(in_1, in_2) ;
  }
}
```

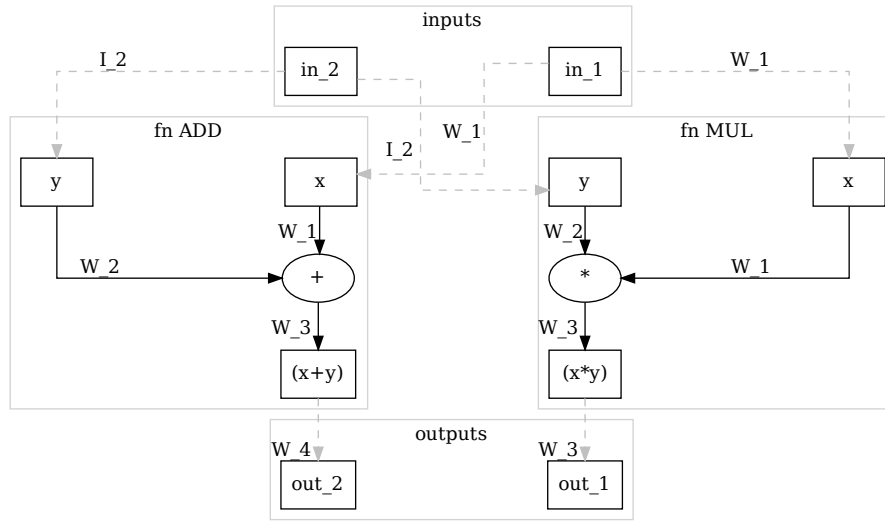
To compile it into an algebraic circuit we start with an empty circuit and evaluate function `main`, which is the only function in this statement.

We draw an inputs subgraph containing box-nodes for every argument of the function and an outputs subgraph containing boxe-nodes for every factor in the return value. Since all of these nodes represent variables of the `field` type, we don't have to add any type constraints to the circuit.

We check the validity of every expression on the right side of every `<==` operator including a type check. In our case every variable is of `field` type and hence the types match the types of the `MUL` as well as the `ADD` function and the type of the left sides of `<==`.


We evaluate the expressions on the right side of every `<==` operator inductively, replacing every occurence of a function by a subgraph that represents its associated circuit.

According to PAPER every occurence of the `public` specifier overwrites the `associate private` default value. Using the appropriate edge labels we get:



Any real world compiler might process its associated high level language in a similar way, replacing functions, or gadgets by predefined associated circuits. This process is then often followed by various optimization steps that try to reduce the number of constraints as much as possible.

In PAPER, we optimize this circuit by collapsing all box nodes that are directly connected to other box nodes, adhering to the rule that a variables `public` specifier overwrites any `private` specifier. Reindexing edge labels we get the following circuit as our pen and pencil compiler output:



Example 128 (3-factorization). Consider our 3-factorization problem from example XXX and the associated circuit $C_{3, \text{fac_zk}}(\mathbb{F}_{13})$ we provided in example XXX. To understand the process of replacing high level functions by their associated circuits, inductively, we want define a PAPER statement, that we brain compile into an algebraic circuit equivalent to $C_{3, \text{fac_zk}}(\mathbb{F}_{13})$. We write

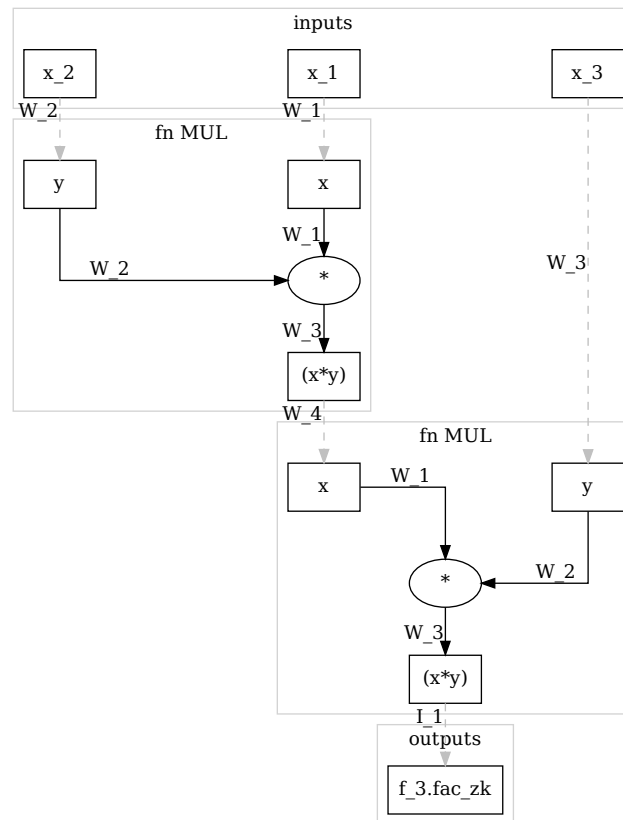
```
statement 3_fac_zk {F:F_13} {
  fn main(x_1 : F, x_2 : F, x_3 : F) -> (pub 3_fac_zk : F) {
    f_3.fac_zk <== MUL( MUL( x_1 , x_2 ) , x_3 ) ;
  }
}
```

Using PAPER, we start with an empty circuit and then add 3 input nodes to the input subgraph as well as 1 output node to the output subgraph. All these nodes are decorated with the associated


variable names. Since all of these nodes represent variables of the `field` type, we don't have to add any type constraints to the circuit.

We check the validity of every expression on the right side of the single `<==` operator including a type check.

We evaluate the expressions on the right side of every `<==` operator inductively. We have two nested multiplication functions and we replace them by the associated multiplication circuits, starting with the most outer function. We get:




In a final optimization step we collapse all box nodes directly connected to other box nodes, adhering to the rule that a variables `public` specifier overwrites any `private` specifier. Reindexing edge labels we get the following circuit:



The Subtraction Constraints System By definition, algebraic circuits only contain addition and multiplication gates and it follows that there is no single gate for field subtraction, despite the fact that subtraction is a native operation in every field.

High level languages and their associated circuit compilers therefore need another way to deal with subtraction. To see how this can be achieved, recall that subtraction is defined by addition with the additive inverse and that the inverse can be computed efficiently by multiplication with -1 . A circuit for field subtraction is therefore given by



Using the general method from XXX, the circuits associated rank-1 constraint system is given by:

$$(S_1 + (-1) \cdot S_2) \cdot 1 = S_3 \quad (7.3)$$

Any valid assignment $\{S_1, S_2, S_3\}$ to this circuit therefore enforces the value S_3 to be the difference $S_1 - S_2$.

Real world compiler usually provide a gadget or a function to abstract over this circuit, such that programmers can use subtraction as if it were native to circuits. In PAPER we define the following subtraction function that compiles to the previous circuit:

```

fn SUB(x : F, y : F) -> (SUB(x,y) : F) {
  constant c : F = -1 ;
  SUB <== ADD(x , MUL( y , c ) ) ;
}

```

In the setup phase of a statement we compile every occurrence of the SUB function into an instance of its associated subtraction circuit and edge labels are generated according to the rules from XXX.


The Inversion Constraint System By definition, algebraic circuits only contain addition and multiplication gates and it follows that there is no single gate for field inversion, despite the fact that inversion is a native operation in every field.

If the underlying field is a prime field, one approach would be to use Fermat's little theorem XXX to compute the multiplicative inverse inside the circuit. To see how this works let \mathbb{F}_p be the prime field. The multiplicative inverse x^{-1} of a field element $x \in \mathbb{F}$ with $x \neq 0$ is then given by $x^{-1} = x^{p-2}$ and computing x^{p-2} in the circuit therefore computes the multiplicative inverse.

Unfortunately, real world primes p are large and computing x^{p-2} by repeated multiplication of x with itself is infeasible. A double and multiply approach as described in XXX is faster as it computes the power in roughly $\log_2(p)$ steps, but still adds a lot of constraints to the circuit.

Computing inverses in the circuit makes no use of the fact, that inversion is a native operation in any field. A more constraints friendly approach is therefore to compute the multiplicative inverse outside of the circuit and then only enforce correctness of the computation in the circuit.

To understand how this can be achieved, observe that a field element $y \in \mathbb{F}$ is the multiplicative inverse of a field element $x \in \mathbb{F}$, if and only if $x \cdot y = 1$ in \mathbb{F} . We can use this and define a circuit, that has two inputs x and y and enforces $x \cdot y = 1$. It is then guaranteed that y is the multiplicative inverse of x . The price we pay is that we can not compute y by circuit execution, but auxillary data is needed to tell any proofer which value of y is needed for a valid circuit assignment. The following circuit defines the constraint



Using the general method from XXX, the circuit is transformed into the following rank-1 constraint system:

$$S_1 \cdot S_2 = 1 \quad (7.4)$$

Any valid assignment $\{S_1, S_2\}$ to this circuit enforces that S_2 is the multiplicative inverse of S_1 and since there is no field element S_2 , such that $0 \cdot S_2 = 1$, it also handles the fact, that the multiplicative inverse of 0 is not defined in any field.

Real world compiler usually provide a gadget or a function to abstract over this circuit and those functions compute the inverse x^{-1} as part of their witness generation process. Programers then don't have to care about providin the inverse as auxiliry data to the circuit. In PAPER we define the following inversion function that compiles to the previous circuit:

```

fn INV(x : F, y : F) -> (x_inv : F) {
  constant c : F = 1 ;
  c <== MUL( x , y ) ;
  x_inv <== y ;
}


```

As we see, this functions takes two inputs, the field value and its inverse. It therefore does not handle the computation of the inverse by itself. This is to keep PAPER as simple as possible.

In the setup phase we compile every occurence of the INV function into an instance of the inversion circuit XXX and edge labels are generated according to the rules from XXX.

The Division Constraint System By definition, algebraic circuits only contain addition and multiplication gates and it follows that there is no single gate for field division, despite the fact that division is a native operation in every field.

Implementing division as a circuit, we use the fact that division is multiplication with the multiplicative inverse. We therefore define division as a circuit using the inversion circuit and constraint system from the prvious paragraph. Expensive inversion is computed outside of the circuit and then provided as circuit input. We get



Using the method from XXX, we transform this circuit into the following rank-1 constraint system:

$$S_2 \cdot S_3 = 1$$

$$S_1 \cdot S_3 = S_4$$

Any valid assignment $\{S_1, S_2, S_3, S_4\}$ to this circuit enforces S_4 to be the field division of S_1 by S_2 . It handles the fact, that division by 0 is not defined, since there is no valid assignment in case $S_2 = 0$.

In PAPER we define the following division function that compiles to the previous circuit:

```
fn DIV(x : F, y : F, y_inv : F) -> (DIV : F) {
  DIV <== MUL( x , INV( y, y_inv ) ) ;
}
```

In the setup phase we compile every occurrence of the binary INV operator into an instance of the inversion circuit.

Exercise 44. Let F be the field \mathbb{F}_5 of modular 5 arithmetics from example XXX. Brain compile the following PAPER statement into an algebraic circuit:

```
statement STUPID_CIRC {F: F_5} {
  fn foo(in_1 : F, in_2 : F)->(out_1 : F, out_2 : F,){
    constant c_1 : F = 3 ;
    out_1<== ADD( MUL( c_1 , in_1 ) , in_1 ) ;
    out_2<== INV( c_1 , in_2 ) ;
  } ;

  fn main(in_1 : F, in_2 : F)->(out_1 : F, out_2 : TYPE_2){
    constant (c_1,c_2) : (F,F) = (3,2) ;
    (out_1,out_2) <== foo(in_1, in_2) ;
  } ;
}
```

Exercise 45. Consider the tiny-jubjub curve from example XXX and its associated circuit XXX. Write a statement in PAPER that brain compiles the statement into a circuit equivalent to the one derived in XXX, assuming that curve points are instances and every other assignment is a witness.

Exercise 46. Let $F = \mathbb{F}_{13}$ be the modular 13 prime field and $x \in F$ some field element. Define a statement in PAPER, such that given instance x a field element $y \in F$ is a witness for the statement, if and only if y is the square root of x .

Brain compile the statement into a circuit and derive its associated rank-1 constraint system. Consider the instance $x = 9$ and compute a constructive proof for the statement.

The Boolean Type


Booleans are a classical primitive type, implemented by virtually every higher programming language. It is therefore of importance to implement booleans in circuits. One of the most common ways to do this is by interpreting the additive and multiplicative neutral element $\{0, 1\} \subset \mathbb{F}$ as the two boolean values, such that 0 represents *false* and 1 represents *true*. Boolean operators like *and*, *or*, or *xor* are then expressible as algebraic computations inside \mathbb{F} .

Representing booleans this way is convenient because the elements 0 and 1 are defined in any field. The representation is therefore independent of the actual field in consideration.

To fix Boolean algebra notation we write 0 to represent *false* and 1 to represent *true* and we write \wedge to represent the boolean AND as well as \vee to represent the boolean OR operator. The boolean NOT operator is written as \neg .

The Boolean Constraint System To represent booleans by the additive and multiplicative neutral elements of a field, a constraint is required to actually enforce variables of boolean type to be either 1 or 0. In fact many of the following circuits that represent boolean functions, are only correct under the assumption that their input variables are constrained to be either 0 or 1. Not constraining boolean variables is a common issue in circuit design.

In order to constrain an arbitrary field element $x \in \mathbb{F}$ to be 1 or 0, the key observation is that the equation $x \cdot (1 - x) = 0$ has only two solutions 0 and 1 in any field. Implementing this equation as a circuit therefore generates the correct constraint:



Using the method from XXX, we transform this circuit into the following rank-1 constraint system:

$$S_1 \cdot (1 - S_1) = 0$$

Any valid assignment $\{S_1\}$ to this circuit enforces S_1 to be either 0 or 1.

Some real world circuit compilers like ZOKRATES or BELLMAN are typed, while others like circom are not. However all of them have their way of dealing with the binary constraint. In PAPER we define the following boolean type that compiles to the previous circuit:

```


type BOOL(b : BOOL) -> (x : F) {
  constant c1 : F = 0 ;
  constant c2 : F = 1 ;
  constant c3 : F = -1 ;
  c1 <== MUL( x , ADD( c2 , MUL( x , c3 ) ) ) ;
  x <== b ;
}

```

In the setup phase of a statement we compile every occurrence of a variable of boolean type into an instance of its associated boolean circuit.

The AND operator constraint system Given two field elements b_1 and b_2 from \mathbb{F} that are constrained to represent boolean variables, we want to find a circuit that computes the logical *and* operator $AND(b_1, b_2)$ as well as its associated RICS, that enforces $b_1, b_2, AND(b_1, b_2)$ to satisfy the constraint system if and only if $b_1 \wedge b_2 = AND(b_1, b_2)$ holds true.

The key insight here is that given three boolean constraint variables b_1, b_2 and b_3 , the equation $b_1 \cdot b_2 = b_3$ is satisfied in \mathbb{F} if and only if the equation $b_1 \wedge b_2 = b_3$ is satisfied in boolean algebra. The logical operator \wedge is therefore implementable in \mathbb{F} by field multiplication of its arguments and the following circuit computes the \wedge operator in \mathbb{F} , assuming all inputs are restricted to be 0 or 1:



The associated rank-1 constraint system can be deduced from the general process XXX and consists of the following constraint

$$S_1 \cdot S_2 = S_3 \quad (7.5)$$


Common circuit languages typically provide a gadget or a function to abstract over this circuit, such that programers can use the \wedge operator without caring about the associated circuit. In PAPER we define the following function that compiles to the \wedge -operator's circuit:

```
fn AND (b_1 : BOOL, b_2 : BOOL) -> AND (b_1, b_2) : BOOL{
  AND (b_1, b_2) <== MUL ( b_1 , b_2 ) ;
}
```

In the setup phase of a statement we compile every occurrence of the AND function into an instance of its associated \wedge -operator's circuit.

The OR operator constraint system Given two field elements b_1 and b_2 from \mathbb{F} that are constrained to represent boolean variables, we want to find a circuit that computes the logical *or* operator $OR(b_1, b_2)$ as well as its associated RICS, that enforces $b_1, b_2, OR(b_1, b_2)$ to satisfy the constraint system if and only if $b_1 \vee b_2 = OR(b_1, b_2)$ holds true.

Assuming that three variables b_1, b_2 and b_3 are boolean constraint, the equation $b_1 + b_2 - b_1 \cdot b_2 = b_3$ is satisfied in \mathbb{F} if and only if the equation $b_1 \vee b_2 = b_3$ is satisfied in boolean algebra. The logical operator \vee is therefore implementable in \mathbb{F} by the following circuit, assuming all inputs are restricted to be 0 or 1:



The associated rank-1 constraint system can be deduced from the general process XXX and consists of the following constraint

$$S_1 \cdot S_2 = S_3$$

$$(S_1 + S_2 - S_3) \cdot 1 = S_4$$

Common circuit languages typically provide a gadget or a function to abstract over this circuit, such that programers can use the \vee operator without caring about the associated circuit. In PAPER we define the following function that compiles to the \vee -operator's circuit:

```

fn OR(b_1 : BOOL, b_2 : BOOL) -> OR(b_1,b_2) : BOOL{
  constant c1 : F = -1 ;
  OR(b_1,b_2) <== ADD(ADD(b_1,b_2),MUL(c1,MUL(b_1,b_2))) ;
}

```

In the setup phase of a statement we compile every occurence of the OR function into an instance of its associated \vee -operator's circuit.


Exercise 47. Let \mathbb{F} be a finite field and let b_1 as well as b_2 two boolean constraint variables from \mathbb{F} . Show that the equation $OR(b_1, b_2) = 1 - (1 - b_1) \cdot (1 - b_2)$ holds true.

Use this equation to derive an algebraic circuit with ingoing variables b_1 and b_2 and outgoing variable $OR(b_1, b_2)$, such that b_1 and b_2 are boolean constraint and the circuit has a valid assignment, if and only if $OR(b_1, b_2) = b_1 \vee b_2$.

Use the technique from XXX to transform this circuit into a rank-1 constraint system and find its full solution set. Define a PAPER function that brain compiles into the circuit.

The NOT operator constraint system Given a field element b from \mathbb{F} that is constrained to represent a boolean variable, we want to find a circuit that computes the logical NOT operator $NOT(b)$ as well as its associated R1CS, that enforces $b, NOT(b)$ to satisfy the constraint system if and only if $\neg b = NOT(b)$ holds true.

Assuming that two variables b_1 and b_2 are boolean constraint, the equation $(1 - b_1) = b_2$ is satisfied in \mathbb{F} if and only if the equation $\neg b_1 = b_2$ is satisfied in boolean algebra. The logical operator \neg is therefore implementable in \mathbb{F} by the following circuit, assuming all inputs are restricted to be 0 or 1:



The associated rank-1 constraint system can be deduced from the general process XXX and consists of the following constraint

$$(1 - S_1) \cdot 1 = S_2$$

Common circuit languages typically provide a gadget or a function to abstract over this circuit, such that programers can use the \neg operator without caring about the associated circuit. In PAPER we define the following function that compiles to the \neg -operator's circuit:

```
fn NOT(b : BOOL -> NOT(b) : BOOL{
  constant c1 = 1 ;
  constant c2 = -1 ;
  NOT(b_1) <== ADD( c1 , MUL( c2 , b) ) ;
}
```

In the setup phase of a statement we compile every occurrence of the NOT function into an instance of its associated \neg -operator's circuit.

Exercise 48. Let \mathbb{F} be a finite field. Derive the algebraic circuit and associated rank-1 constraint system for the following operators: NOR, XOR, NAND, EQU.

Modularity As we have seen in XXX and XXX, both algebraic circuits and R1CS have a modularity property and as we have seen in this section, all basic boolean functions are expressible in circuits. Combining those two properties, show that it is possible to express arbitrary boolean functions as algebraic circuits.

This shows that the expressiveness of algebraic circuits and therefore rank-1 constraint systems is as general as the expressiveness of boolean circuits. In important implication is that the languages $L_{R1CS-SAT}$ and $L_{Circuit-SAT}$ as defined in XXX, are as general as the famous language L_{3-SAT} , which is known to be \mathcal{NP} -complete.

Example 129. To give an example of how a compiler might construct complex boolean expressions in algebraic circuits from simple one and how to derive their associated rank-1 constraint systems, lets look at the following PAPER statement:

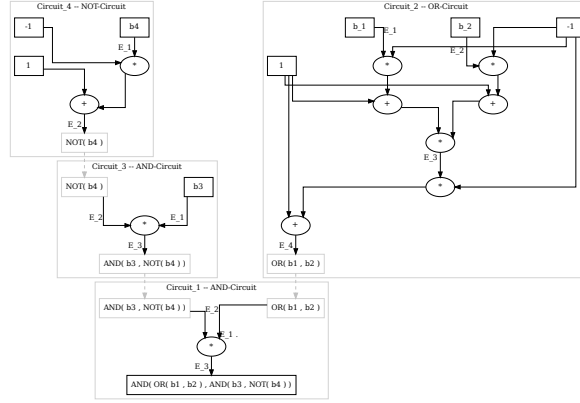
```
statement BOOLEAN_STAT {F: F_p} {
  fn main(b_1:BOOL,b_2:BOOL,b_3:BOOL,b_4:BOOL )-> pub b_5:BOOL {
    b_5 <== AND( OR( b_1 , b_2) , AND( b_3 , NOT( b_4) ) ) ;
  } ;
}
```

The code describes a circuit, that takes four private inputs b_1, b_2, b_3 and b_4 of boolean type and computes a public output b_5 , such that the following boolean expression holds true:

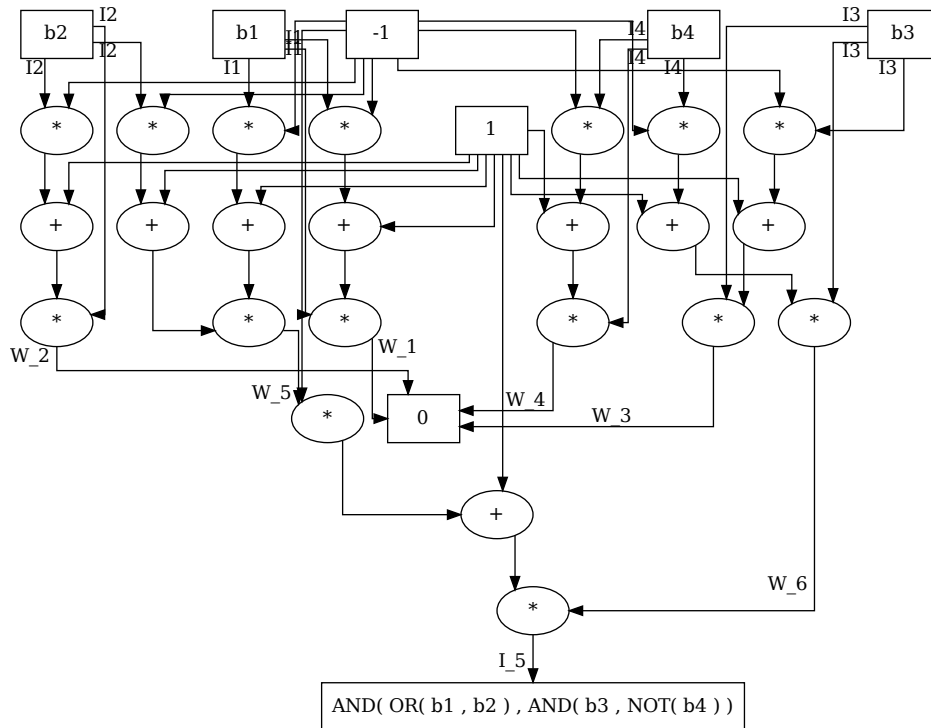
$$(b_1 \vee b_2) \wedge (b_3 \wedge \neg b_4) = b_5$$

During a setup-phase, a circuit compilers then transforms this high level language statement into a circuit and associated rank-1 constraints systems and hence defines a language $L_{BOOLEAN_STAT}$.

To see how this might be achieved, we use PAPER as an example to execute the setup-phase and compile `BOOLEAN_STAT` into a circuit. Taking the definition of the boolean constraint `XXX` as well as the definitions of the appropriate boolean operators into account, we get the following circuit:



Simple optimization then collapses all box-nodes that are directly linked and all box nodes that represent the same constants. After relabeling the edges the following circuit represents the circuit associated to the `BOOLEAN_STAT` statement:



Given some public input I_1 from \mathbb{F}_{13} a valid assignments to this circuits consists of private inputs W_1, W_2, W_3, W_4 from \mathbb{F}_{13} , such that the equation $I_1 = (W_1 \vee W_2) \wedge (W_3 \wedge \neg W_4)$ holds true. In addition a valid assignment also has to contain private inputs W_5, W_6, W_7, W_8, W_9 and W_{10} ,

which can be derived from circuit execution. The inputs W_5, \dots, W_8 ensure that the first four private inputs are either 0 or 1 but not any other field element and the others enforce the boolean operations in the expression.

To compute the associated R1CS we can use the general method from XXX and look at every labeled outgoing edge not coming from a source node. Declaring the edges coming from input nodes as well as the edge going to the single output node as public and every other edge as private input. In this case we get:

$$\begin{aligned}
 W_5 : W_1 \cdot (1 - W_1) &= 0 && \text{boolean constraints} \\
 W_6 : W_2 \cdot (1 - W_2) &= 0 \\
 W_7 : W_3 \cdot (1 - W_3) &= 0 \\
 W_8 : W_4 \cdot (1 - W_4) &= 0 \\
 W_9 : W_1 \cdot W_2 &= W_9 && \text{first OR-operator constraint} \\
 W_{10} : W_3 \cdot (1 - W_4) &= W_{10} && \text{AND(.,NOT(.))-operator constraints} \\
 I_1 : (W_1 + W_2 - W_9) \cdot W_{10} &= I_1 && \text{AND-operator constraints}
 \end{aligned}$$

The reason why this R1CS only contains a single constraint for the multiplication gate in the OR-circuit, while the general definition XXX requires two constraints, is that the second constraint in XXX only appears since the final addition gate is connected to an output node. In this case however the final addition gate from the OR-circuit is enforced in the left factor of the I_1 constraint. Something similar holds true for the negation circuit.

During a proofer-phase, some public instance I_5 must be given. To compute a constructive proof for the statement of the associated languages with respect to instance I_5 , a proofer has to find four boolean values W_1, W_2, W_3 and W_4 , such that

$$(W_1 \vee W_2) \wedge (W_3 \wedge \neg W_4) = I_5$$

holds true. In our case neither the circuit, nor the PAPER statement specifies how to find those values and it is a problem that any proofer has to solve outside of the circuit. This might or might not be true for other problems, too. In any case once the proofer found those values, they can execute the circuit to find a valid assignment.

To give a concrete example let $I_1 = 1$ and assume $W_1 = 1, W_2 = 0, W_3 = 1$ and $W_4 = 0$. Since $(1 \vee 0) \wedge (1 \wedge \neg 0) = 1$ those values satisfy the problem and we can use them to execute the circuit. We get

$$\begin{aligned}
 W_5 &= W_1 \cdot (1 - W_1) = 0 \\
 W_6 &= W_2 \cdot (1 - W_2) = 0 \\
 W_7 &= W_3 \cdot (1 - W_3) = 0 \\
 W_8 &= W_4 \cdot (1 - W_4) = 0 \\
 W_9 &= W_1 \cdot W_2 = 0 \\
 W_{10} &= W_3 \cdot (1 - W_4) = 1 \\
 I_1 &= (W_1 + W_2 - W_9) \cdot W_{10} = 1
 \end{aligned}$$

A constructive proof of knowledge of a witness for instance $I_1 = 1$ is therefore given by the tuple $P = (W_5, W_6, W_7, W_8, W_9, W_{10}) = (0, 0, 0, 0, 0, 1)$.

Arrays

The `array` type represents a fixed size collection of elements of equal type, each selectable by one or more indices that can be computed at run time during program execution.

Arrays are a classical type, implemented by many higher programming language that compile to circuits or rank-1 constraints systems, however most high level circuit languages supports *static* arrays, i.e. arrays whose length is known at compile time, only.

The most common way to compile arrays to circuits is to transform any array of a given type t and size N into N circuit variables of type N . Arrays are therefore syntactic sugar, that the compiler transforms into input nodes, much like any other variable. In PAPER we define the following array type:

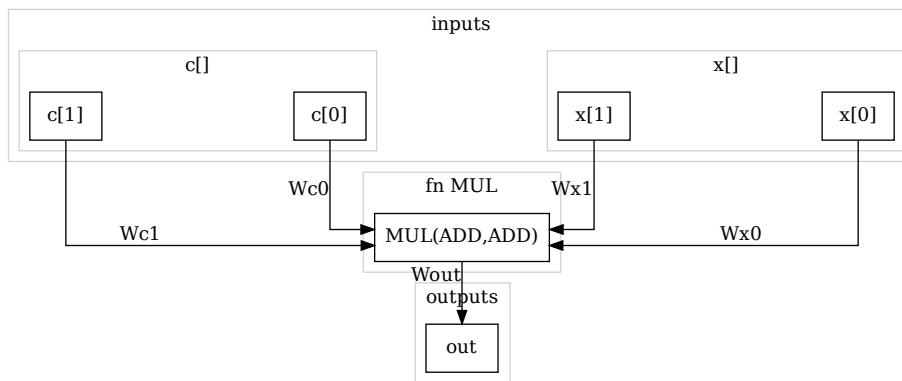
```
type <Name>: <Type>[N : unsigned] -> (Type,...) {
  return (<Name>[0],...)
}
```

In the setup phase of a statement we compile every occurrence of an array of size N that contains elements of type `Type` into N variables of type `Type`.

Example 130. To give an intuition of how a real world compiler might transform arrays into circuit variables, consider the following PAPER statement:

```
statement ARRAY_TYPE {F: F_5} {
  fn main(x: F[2]) -> F {
    let constant c: F[2] = [2,4] ;
    let out:F <== MUL(ADD(x[1],c[0]),ADD(x[0],c[1])) ;
    return out ;
  } ;
}
```

During a setup-phase, a circuit compiler might then replace any occurrence of the array type by a tuple of variables of the underlying type and then use those variables in the circuit synthesis process instead. To see how this can be achieved, we use PAPER as an example. Abtracting over the subcircuit of the computation, we get the following circuit:



The Unsigned Integer Type

Unsigned integers of size N , where N is usually a power of two represents non negative integers in the range $0 \dots 2^N - 1$. They have a notion of addition, subtraction and multiplication, defined by modular 2^N arithmetics. If some N is given, we write uN for the associated type.

The uN Constraints System Many high level circuit languages define the the various uN types as arrays of size N, where each element is of boolean type. This is similar to their representation on common compuer hardware and allows for efficient and straight forward definition of common operators, like the various shift, or logical operators.

If some unsigned integer N is known at compile time, in PAPER we define the following uN type:

```
type uN -> BOOL[N] {
  let base2 : BOOL[N] <== BASE_2 (uN) ;
  return base2 ;
}
```

To enforce an N -tuple of field elements (b_0, \dots, b_{N-1}) to represent an element of type uN we therefore need N boolean constraints

$$\begin{aligned} S_0 \cdot (1 - S_0) &= 0 \\ S_1 \cdot (1 - S_1) &= 0 \\ &\dots \\ S_{N-1} \cdot (1 - S_{N-1}) &= 0 \end{aligned}$$

In the setup phase of a statement we compile every occurence of the uN type by a size N array of boolean type. During a proofer phase actual elements of the uN type are first transformed into binary representation and then this binary representation is assigned to the boolean array that represents the uN type.

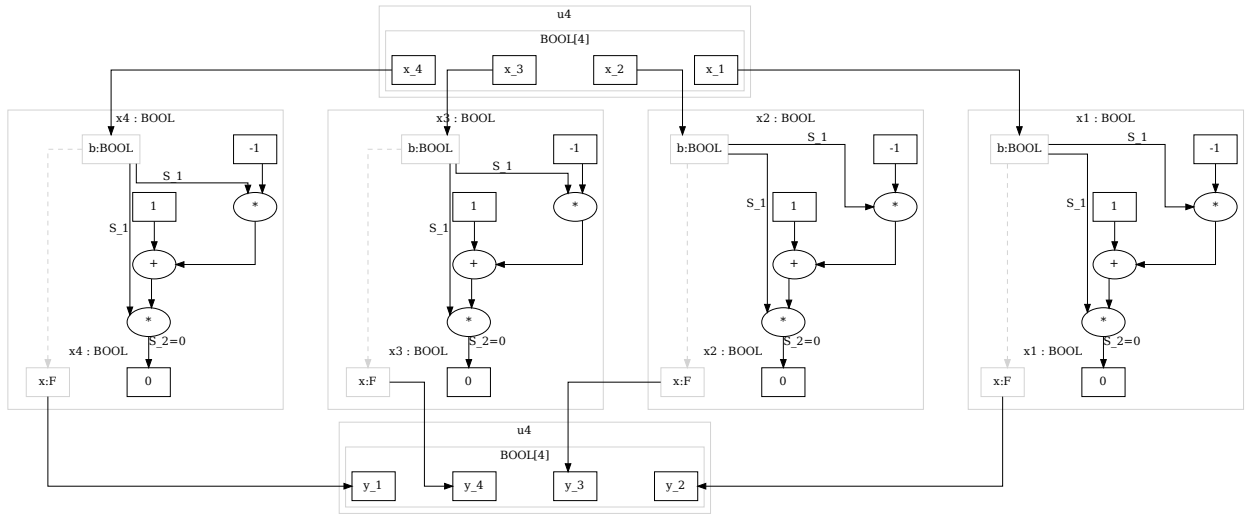
Remark 4. Representing the uN type as boolean arrays is conceptually clean and works over generic base fields. However representing unsigned integers in this way requires a lot of space as every bit is represented as a field element and if the base field is large, those field elements require considerable space in hardware.

It should be noted that in some cases there is another more space and constraints efficient approach to represent unsigned integers, that can be used whenever the underlying base field is sufficiently large. To understand this, recall that addition and multiplication in a prime field \mathbb{F}_p is equal to addition and multiplication of integers, as long as the sum or the product does not exceed the modulus p . It is therefore possible to represent the uN type inside the basefield type, whenever N is small enough. In this case however care has to be taken to never overflow the modulus. It is also important to make sure that in subtraction the subtrahend is never larger then the minuend.

Example 131. To give an intuition of how a real world compiler might transform unsigned integers into circuit variables, consider the following PAPER statement:

```
statement RING_SHIFT{F: F_p, N=4} {
  fn main(x: uN)-> uN {
    let y:uN <== [x[1],x[2],x[3],x[0]] ;
    return y ;
  } ;
}
```

During a setup-phase, a circuit compiler might then replace any occurence of the uN type by N variables of boolean type. Using the definition of booleans each of these variables is then transformed into the field type and a boolean constraints sstem. To see how this can be achieved, we use PAPER as an example and get the following circuit:



During a proofer phase function `main` is called with an actual input of `u4` type, say `x=14`. The high level language then has to transform the decimal value 14 into its 4-bit binary representation $14_2 = (0, 1, 1, 1)$ outside of the circuit. Then the array of field values $x[4] = [0, 1, 1, 1]$ is used as input to the circuit. Since all 4 field elements are either 0 or 1 the four boolean constraints are satisfiable and the output is an array of the four field elements $[1, 1, 1, 0]$, which represents the `u4` element 7.

The Unigned Integer Operators Since elements of `uN` type are represented as boolean arrays, shift operators are implement in circuits simply by rewiring the boolean input variables to the output variables accordingly.

Logical operators, like AND, OR, or NOT are defined on the `uN` type by invoking the appropriate boolean operators bitwise to every bit in the boolean array that represents the `uN` element.

Addition and multiplication can be represented similar to how machines represent those operations. Addition can be implemented by first defining the *full adder* circuit and then combining N of this these circuits into a circuit that add to elements from the `uN` type.

Exercise 49. Let $F = \mathbb{F}_{13}$ and $N=4$ be fixed. Define circuits and associated R1CS for the left and right bishift operators $x \ll 2$ as well as $x \gg 2$ that operate on the `uN` type. Execute the associated circuit for $x : u4 = 11$.

Exercise 50. Let $F = \mathbb{F}_{13}$ and $N=2$ be fixed. Define a circuit and associated R1CS for the addition operator $\text{ADD} : F \times F \rightarrow F$. Execute the associated circuit to compute $\text{ADD}(2, 7)$.

Exercise 51. Brain compile the following PAPER code into a circuit and derive the associated R1CS.

```
statement MASK_MERGE {F:F_5, N=4} {
  fn main(pub a : uN, pub b : uN) -> F {
    let constant mask : uN = 10 ;
    let r : uN <== XOR(a, AND(XOR(a,b), mask)) ;
    return r ;
  }
}
```

Let L_{mask_merge} be the language defined by the circuit. Provide a constructive knowledge proof in L_{mask_merge} for the instance $I = (I_a, I_b) = (14, 7)$.

7.2.2 Control Flow

Most programming languages of the imperative or functional style have some notion of basic control structures to direct the order in which instructions are evaluated. Contemporary circuit compilers usually provide a single thread of execution and provide basic flow constructs that implement control flow in circuits.

The Conditional Assignment

Writing high level code that compiles to circuits, it is often necessary to have a way for conditional assignment of values or computational output to variables.

One way to realize this in many programming languages is in terms of the conditional ternary assignment operator $?:$ that branches the control flow of a program according to some condition and then assigns the output of the computed branch to some variable.

```
variable = condition ? value_if_true : value_if_false
```

In this description `condition` is a boolean expression and `value_if_true` as well as `value_if_false` are expressions that evaluate to the same type as `variable`.

In programming languages like Rust another way to write the conditional assignment operator that is more familiar to many programmers is given by


```
variable = if condition then {  
    value_if_true  
} else {  
    value_if_false  
}
```

In most programming languages it is a key property of the ternary assignment operator that the expression `value_if_true` is only evaluated if `condition` evaluates to true and the expression `value_if_false` is only evaluated if `condition` evaluates to false. In fact computer programs would turn out to be very inefficient if the ternary operator would evaluate both expressions regardless of the value of `condition`.

A simple way to implement conditional assignment operator as a circuit can be achieved, if the requirement that only one branch of the conditional operator is executed is dropped. To see that let b , c and d be field elements, such that b is boolean constraint. In this case the following equation enforces a field element x to be the result of the conditional assignment operator:

$$x = b \cdot c + (1 - b) \cdot d \quad (7.6)$$

Expressing this equation in terms of the addition and multiplication operators from XXX, we can flatten it into the following algebraic circuit:



Note that in order to compute a valid assignment to this circuit, both S_2 as well as S_4 are necessary. If the inputs to the nodes c and d are circuits themselves, both circuits need valid assignments and therefore have to be executed. As a consequence this implementation of the conditional assignment operator has to execute all branches of all circuits, which is very different from the execution of common computer programs and contributes to the increased computational effort any prover has to invest, in contrast to the execution in other programming models.

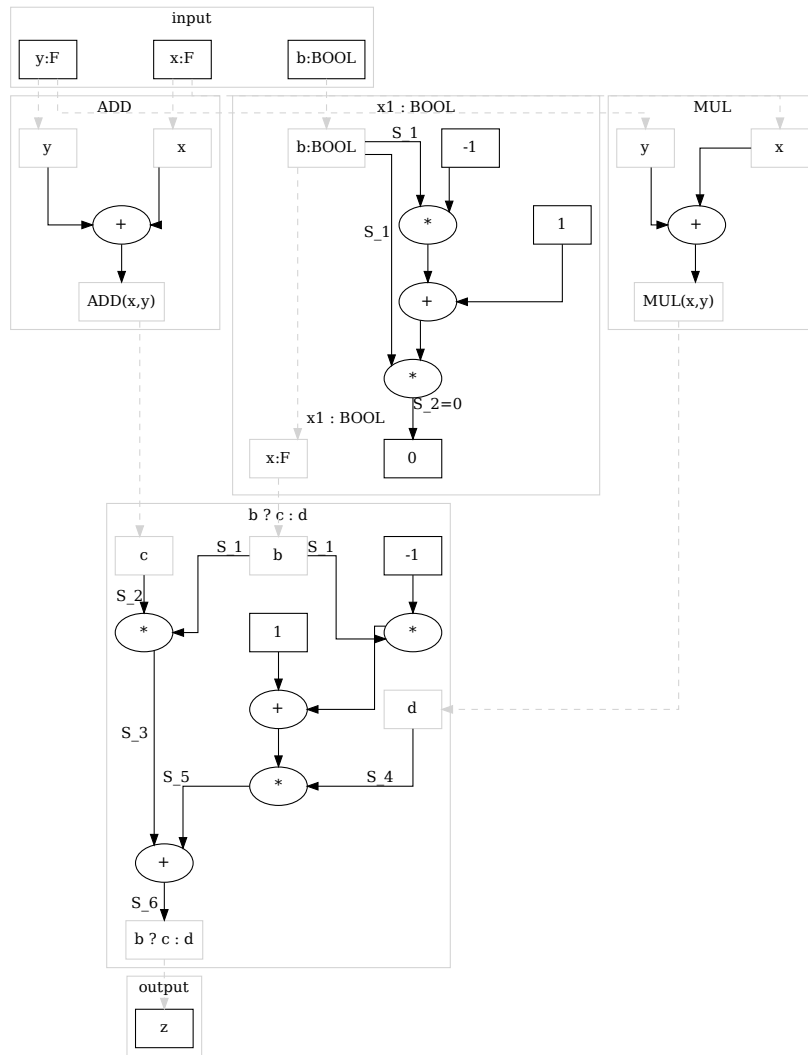
We can use the general technique from XXX to derive the associated rank-1 constraint system of the conditional assignment operator. We get

$$\begin{aligned} S_1 \cdot S_2 &= S_3 \\ (1 - S_1) \cdot S_4 &= S_5 \\ (S_3 + S_5) \cdot 1 &= S_6 \end{aligned}$$

Example 132. To give an intuition of how a real world circuit compiler might transform any high level description of the conditional assignment operator into a circuit, consider the following PAPER code:

```
statement CONDITIONAL_OP {F:F_p} {
  fn main(x : F, y : F, b : BOOL) -> F {
    let z : F <== if b then {
      ADD(x,y)
    } else {
      MUL(x,y)
    } ;
    return z ;
  }
}
```

Brain compiling this code into a circuit, we first draw box nodes for all input and output variables and then transform the boolean type into the field type together with its associated constraint. Then we evaluate the assignments to the output variables. Since the conditional assignment operator is the top level function, we draw its circuit and then draw the circuits for both conditional expressions. We get:



Loops

In many programming languages various loop control structures are defined that allow developers to execute expressions with a specified number of repetitions or arguments. In particular it is often possible to implement unbounded loops like the

```
while true do { }
```

structure, or loop structure, where the number of executions depends on execution inputs and is therefore unknown at compile time.

In contrast to this it should be noted that algebraic circuits and rank-1 constraints systems are not general enough to express arbitrary computation, but bounded computation only. As a consequence it is not possible to implement unbounded loops, or loops with bounds that are unknown at compile time in those models. This can be easily seen since circuits are acyclic by definition and implementing an unbounded loop as an acyclic graph requires a circuits of unbounded size.

However circuits are general enough to express bounded loops, where the upper bound on its execution is known at compile time. Those loop can be implemented in circuits by enrolling the loop.

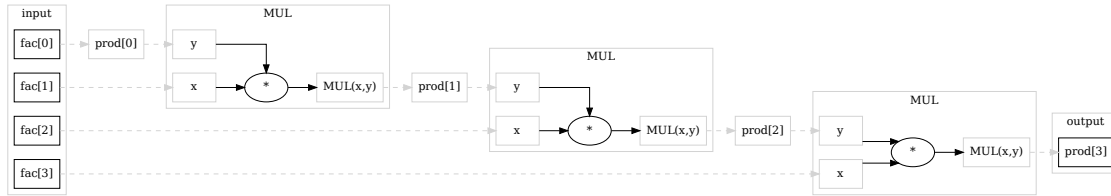
As a consequence, any programming language that compiles to algebraic circuits can only provide loop structures, where the bound is a constant known at compile-time. This implies that loops cannot depend on execution inputs, but on compile time parameters, only.

Example 133. To give an intuition of how a real world circuit compiler might transform any high level description of a bounded `for` loop into a circuit, consider the following PAPER code:

```
statement FOR_LOOP {F:F_p, N: unsigned = 4} {
  fn main(fac : F[N]) -> F {
    let prod[N] : F ;
    prod[0] <== fac[0] ;
    for unsigned i in 1..N do [{
      prod[i] <== MUL(fac[i], prod[i-1]) ;
    }]
    return prod[N] ;
  }
}
```

Note that in a program like this, the loop counter `i` has no expression in the derived circuit. It is pure syntactic sugar, telling the compiler how to unrole the loop.

Brain compiling this code into a circuit, we first draw box nodes for all input and output variables, noting that the loop counter is not represented in the circuit. Since all variables are of `field` type, we don't have to compile any type constraints. Then we evaluate the assignments to the output variables, by unrolling the loop into 3 individual assignment operators. We get:



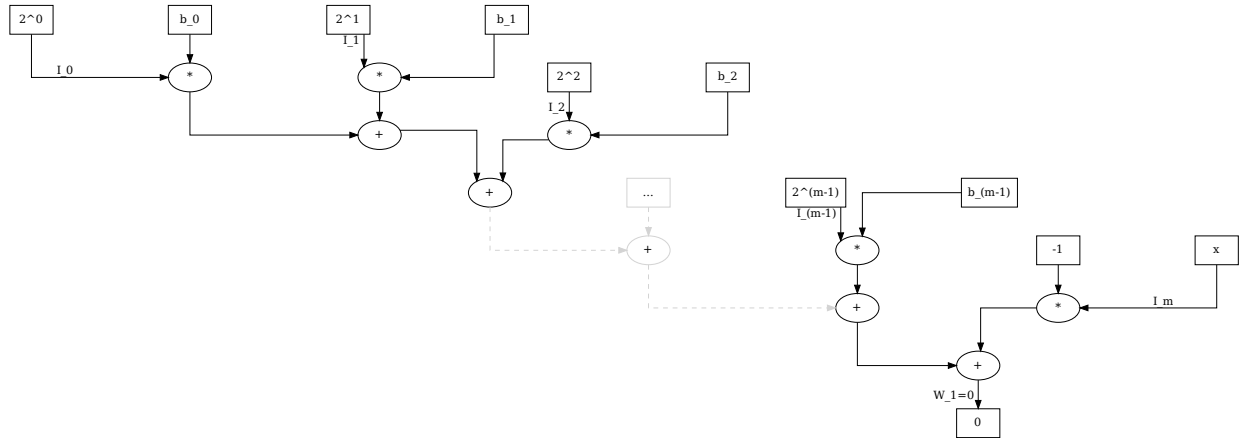
7.2.3 Binary Field Representations

In applications is often necessary to enforce a binary representation of elements from the `field` type. To derive an appropriate circuit over a prime field \mathbb{F}_p , let $m = |p|_2$ be the smallest number of bits necessary to represent the prime modulus p . Then a bitstring $(b_0, \dots, b_{m-1}) \in \{0, 1\}^m$ is a binary representation of a field element $x \in \mathbb{F}_p$, if and only if

$$x = b_0 \cdot 2^0 + b_1 \cdot 2^1 + \dots + b_{m-1} \cdot 2^{m-1}$$

In this expression, addition and exponentiation is considered to be executed in \mathbb{F}_p , which is well defined, since all terms 2^j for $0 \leq j < m$ are elements of \mathbb{F}_p . Note however that in contrast to the binary representation of unsigned integers $n \in \mathbb{N}$, this representation is not unique in general, since the modular p equivalence class might contain more then one binary representative.

Considering that the underlying prime field is fixed and the most significant bit of the prime modulus is m , the following circuit flattens equation XXX, assuming all inputs b_1, \dots, b_m are of boolean type.



Applying the general transformation rule to compute the associated rank-1 constraint systems, we see that we actually only need a single constraint to enforce some binary representation of any field element. We get

$$(S_0 \cdot 2^0 + S_1 \cdot 2^1 + \dots + S_{m-1} \cdot 2^{m-1} - S_m) \cdot 1 = 0$$

Given an array `BOOL[N]` of N boolean constraint field elements and another field element x , the circuit enforces `BOOL[N]` to be one of the binary representations of x . If `BOOL[N]` is not a binary representation of x , no valid assignment and hence no solution to the associated R1CS can exist.

Example 134. Consider the prime field \mathbb{F}_{13} . To compute binary representations of elements from that field, we start with the binary representation of the prime modulus 13, which is $|13|_2 = (1, 0, 1, 1)$ since $13 = 1 \cdot 2^0 + 0 \cdot 2^1 + 1 \cdot 2^2 + 1 \cdot 2^3$. So $m = 4$ and we need up to 4 bits to represent any element $x \in \mathbb{F}_{13}$.

To see that binary representations are not unique in general, consider the element $2 \in \mathbb{F}_{13}$. It has the binary representations $|2|_2 = (0, 1, 0, 0)$ and $|2|_2 = (1, 1, 1, 1)$, since in \mathbb{F}_{13} we have

$$2 = \begin{cases} 0 \cdot 2^0 + 1 \cdot 2^1 + 0 \cdot 2^2 + 0 \cdot 2^3 \\ 1 \cdot 2^0 + 1 \cdot 2^1 + 1 \cdot 2^2 + 1 \cdot 2^3 \end{cases}$$

This is because the unsigned integers 2 and 15 are both in the modular 13 remainder class of 2 and hence are both representatives of 2 in \mathbb{F}_{13} .

To see how circuit XXX works, we want to enforce the binary representation of $7 \in \mathbb{F}_{13}$. Since $m = 4$ we have to enforce a 4-bit representation for 7, which is $(1, 1, 1, 0)$, since $7 = 1 \cdot 2^0 + 1 \cdot 2^1 + 1 \cdot 2^2 + 0 \cdot 2^3$. A valid circuit assignment is therefore given by $(S_0, S_1, S_2, S_3, S_4) = (1, 1, 1, 0, 7)$ and indeed the assignment satisfies the required 5 constraints including the 4 boolean constraints for S_0, \dots, S_3 :

$$\begin{aligned} 1 \cdot (1 - 1) &= 0 & // \text{boolean constraints} \\ 1 \cdot (1 - 1) &= 0 \\ 1 \cdot (1 - 1) &= 0 \\ 0 \cdot (1 - 0) &= 0 \\ (1 + 2 + 4 + 0 - 7) \cdot 1 &= 0 & // \text{binary rep. constraint} \end{aligned}$$

7.2.4 Cryptographic Primitives

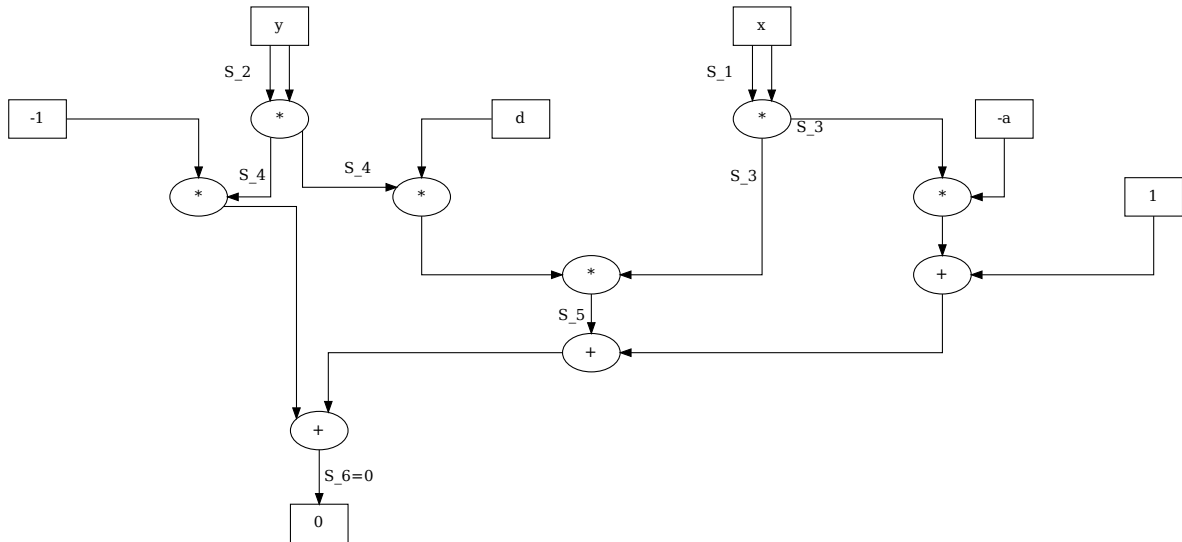
In applications it is often required to do cryptography in a circuit. To do this, basic cryptographic primitives, like hash functions or elliptic curve cryptography needs to be implemented as circuits. In this section we give a few basic examples of how to implement those primitives.

Twisted Edwards curves

Implementing elliptic curve cryptography in circuits, means to implement the field equation as well as the algebraic operations of an elliptic curve as circuits and to do this efficiently the curve must be defined over the same base field as the field that is used in the circuit.

For efficiency reasons it is advantageous to choose an elliptic curve, such that that all required constraints and operations can be implement with as few gates as possible. Twisted Edwards curves are particularly useful for that matter, since their addition law is particularly simple and the same equation can be used for all curve points including the point at infinity. This simplifies the circuit a lot.

Twisted Edwards curves constraints As we have seen in XXX, a twisted Edwards curve over a finite field F is defined as the set of all pairs of points $(x, y) \in \mathbb{F} \times \mathbb{F}$, such that x and y satisfy the equation $a \cdot x^2 + y^2 = 1 + d \cdot x^2 y^2$. As we have seen in example XXX, we can transform this equation into the following circuit:



The circuit enforces the two inputs of `field` type to satisfy the twisted Edwards curve equation and as we know from example XXX, the associated rank-1 constraints system is given by:

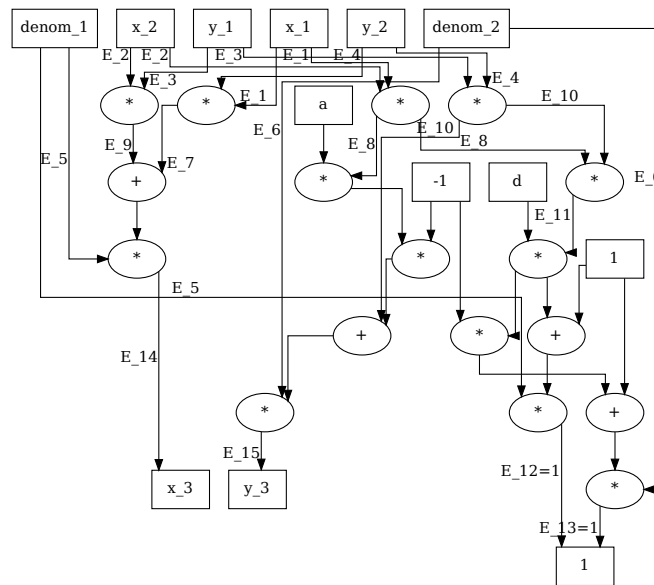
$$\begin{aligned}
 S_1 \cdot S_1 &= S_3 \\
 S_2 \cdot S_2 &= S_4 \\
 (S_4 \cdot 8) \cdot S_3 &= S_5 \\
 (12 \cdot S_4 + S_5 + 10 \cdot S_3 + 1) \cdot 1 &= 0
 \end{aligned}$$

Exercise 52. Write the circuit and associated rank-1 constraint system for a Weierstraß curve of a given field \mathbb{F} .

Twisted Edwards curves addition As we have seen in XXX a major advantage of twisted Edwards curves is the existence of an addition law, that contains no branching and is valid for all curve points. Moreover the neutral element is not "at infinity" but the actual curve poin $(0, 1)$. In fact given two points (x_1, y_1) and (x_2, y_2) on a twisted Edwards curve their sum is defined as

$$(x_3, y_3) = \left(\frac{x_1 y_2 + y_1 x_2}{1 + d \cdot x_1 x_2 y_1 y_2}, \frac{y_1 y_2 - a \cdot x_1 x_2}{1 - d \cdot x_1 x_2 y_1 y_2} \right)$$

We can use the division circuit from XXX to flatten this equation into an algebraic circuit. Inputs to the circuit are then not only the two curve points (x_1, y_1) and (x_2, y_2) but also the two denominators $denum_1 = 1 + d \cdot x_1 x_2 y_1 y_2$ as well as $denum_2 = 1 - d \cdot x_1 x_2 y_1 y_2$, which any prover needs to compute outside of the circuit. We get



Using the general technique from XXX to derive the associated rank-1 constraint system, we get the following result:

$$\begin{aligned} S_1 \cdot S_4 &= S_7 \\ S_1 \cdot S_2 &= S_8 \\ S_2 \cdot S_3 &= S_9 \\ S_3 \cdot S_4 &= S_{10} \\ S_8 \cdot S_{10} &= S_{11} \\ S_5 \cdot (1 + d \cdot S_{11}) &= 1 \\ S_6 \cdot (1 - d \cdot S_{11}) &= 1 \\ S_5 \cdot (S_9 + S_7) &= S_{14} \\ S_6 \cdot (S_{10} - a \cdot S_8) &= S_{15} \end{aligned}$$

Exercise 53. Let \mathbb{F} be a field. Define a circuit that enforces field inversion for a point of a twisted Edwards curve over \mathbb{F} .

Bibliography

Jens Groth. On the size of pairing-based non-interactive arguments. *IACR Cryptol. ePrint Arch.*, 2016:260, 2016. URL <http://eprint.iacr.org/2016/260>.

David Fifield. The equivalence of the computational diffie–hellman and discrete logarithm problems in certain groups, 2012. URL <https://web.stanford.edu/class/cs259c/finalpapers/dlp-cdh.pdf>.

Torben Pryds Pedersen. Non-interactive and information-theoretic secure verifiable secret sharing. In Joan Feigenbaum, editor, *Advances in Cryptology — CRYPTO '91*, pages 129–140, Berlin, Heidelberg, 1992. Springer Berlin Heidelberg. ISBN 978-3-540-46766-3. URL <https://fmouhart.epheme.re/Crypto-1617/TD08.pdf>.

Martin Albrecht, Lorenzo Grassi, Christian Rechberger, Arnab Roy, and Tyge Tiessen. Mimc: Efficient encryption and cryptographic hashing with minimal multiplicative complexity. Cryptology ePrint Archive, Report 2016/492, 2016. <https://ia.cr/2016/492>.